

Detection and Analysis of Disaster-Related Tweets

Daniel Solomon
solomond@mail.tau.ac.il

Gal Ron
galr1@mail.tau.ac.il

Omri Ben-Horin
EMAIL@mail.tau.ac.il

Abstract

[illegible]

1 Introduction

The popular microblogging service Twitter is a fruitful source of user-created content. With hundreds of millions of new tweets every day, Twitter has become a probe to human behavior and opinions from around the globe. The Twitter 'corpus' reflects political and social trends, popular culture, global and local happenings, and more. In addition, tweets are easy to access and aggregate in real-time. Therefore, we experience an increased interest in natural language processing research of Twitter data.

As one of the world's most widely used social networks, Twitter is an effective channel of communication and plays an important role during a crisis or emergency. The live stream of tweets can be used to identify reports and calls for help in emergency situations, such as accidents, violent crimes, natural disasters and terror attacks (which we all refer to as 'disasters' in this paper).

In this work we utilize techniques from the natural language processing pipeline (tokenization, part-of-speech tagging and named-entity recognition) to work on Twitter data, as opposed to traditional corpora, in order to detect and analyze disaster-related tweets.

The Dataset We present our experiments on a dataset of 10,877 tweets¹, labeled to '*disaster-related*' and '*not disaster-related*' with confidence in the range $[0, 1]$. For

example, the following tweet is '*disaster-related*' with confidence 1.

Thunderstorms with little rain expected
in Central California. High fire danger.
#weather #cawx <http://t.co/A5GNzbuSqq>

while the following tweet is '*not disaster-related*' with confidence 0.59.

It's been raining since you left me // Now I'm
drowning in the flood // You see I've always
been a fighter // But without you I give up

Even for one who is not familiar with the Bon Jovi lyrics in the latter tweet, it is clear that the tweet does not refer to a real natural disaster. However, this observation is hard to make examining only the vocabulary used; the latter tweet contains a variety of 'disastrous' words (e.g. raining, drowning, flood, fighter). This example hints that in order to reach meaningful results we may have to examine additional linguistic features of colloquial writing, as well as Twitter-specific features such as hashtags (#), user at-mentions (@), internet links and emoticons.

Our Contribution In this paper we present our work tackling three missions involving disaster-related tweets.

The first mission is *identification* of disaster-related tweets among a variety of tweets. We implemented several classifiers, the best of which achieved **TODO: %** accuracy on the dataset. We note that this method could have easily been adjusted to identify tweets related to themes other than disasters (e.g. politics-related, sports-related, etc.), given the appropriate dataset.

The second mission is binary classification of disaster-related tweets to one of two categories, *subjective tweets* (i.e. tweets that express an emotion) vs. *objective tweets* (such as news reports on disasters). To achieve this we manually tagged 2,410 disaster-related tweets. The motivation behind this task is that objective tweets like informative news reports are likely to be published after

the event had already become clear to emergency services, while subjective tweets may contain invaluable first-person testimonies of ongoing events.

Finally, we extracted named entities to enrich our knowledge on the disaster (mostly location)... **TODO: Omri - short description of method and achievements.**

To demonstrate our framework we aggregated recent tweets from various locations in the US, extracted disaster-related tweets using our classifier, and then used named-entity recognition to discover entities related to ongoing disasters. For example, "Hurricane Harvey" appeared as a top named-entity among recent tweets sent from Houston, TX, which we identified as *disaster-related*.

The code of our project is available at <https://github.com/glrn/nlp-disaster-analysis>.

1.1 Twitter vs. Traditional Corpora

Tweets are limited to 140 characters and are widely used by non-professional writers. Therefore, Tweet datasets have some unique features that differ from traditional corpora (such as WSJ corpus). These features should be addressed when implementing natural language processing techniques.

First, the grammar of tweets is quite different from edited news text. It is common that tweets are written as colloquial sentences in first person where the subject ('I') is omitted, as in: 'see the flames from my window OMG'.

Tweets are also characterized by extensive use of abbreviations and slang specific to social-media (e.g. 'ily' for 'I love you', 'iono' for 'I don't know'). Such abbreviations may squash several parts-of-speech into one token, which poses a challenge to POS tagging.

In addition, due to the colloquial nature of user-created content, it is common that proper words are replaced by phonetically or morphologically similar ones (e.g. 'wtchng' instead of 'watching', 'gr8' instead of 'great'). Users may also use capitalization irregularities, deliberate spelling errors, punctuation irregularities and interjections as a means to express their sentiment, as in the following tweet:

Haha South Tampa is getting flooded hah-
WAIT A SECOND I LIVE IN SOUTH
TAMPA WHAT AM I GONNA DO WHAT AM
I GONNA DO FVCK #flooding

Lastly, tweets may contain a variety of tokens seen mainly in Twitter and other social media, such as: URLs; emoticons; Twitter hashtags, of the form #tagname, which the author may supply to label a tweet; Twitter mentions of the form @user which link to other Twitter

users; and Twitter discourse functions such as RT ("re-tweet"), indicating that a tweet was originally posted by some other Twitter user, or ellipsis dots (...) at the end (or beginning) of a tweet, indicating that the tweet will be continued in a subsequent tweet by the same user. We note that hashtags and at-mentions can also serve as words or phrases within a tweet, as in:

Heard about the #earthquake, stay safe everyone.

Regarding URLs, all links posted in tweets are shortened using Twitter's link service, <http://t.co>, and are converted to a seemingly random 23 characters URL that redirects to the original web address.

2 Analysis Workflow

keywords TODO

- A
- B
- C

TODO: (Gal) Complete this section

3 Classification of Disaster-Related Tweets

In the first part of our work we developed a classifier that identifies *disaster-related* tweets from *not disaster-related* tweets, trained on a dataset of 10,877 labeled tweets (**TODO: note: maybe better to say that the dataset was 2400 tweets, since we looked only at tweets with confident > 0.9**). We experimented with a Naive Bayes (NB), random forest (RF) and support vector machine (SVM) classifiers.

Naive Bayes TODO: (Daniel) short explanation

Support Vector Machine (SVM) TODO: (Daniel) short explanation

Random Forest TODO: (Daniel) short explanation

3.1 Feature Extraction

To train the classifiers we extracted several features for each tweet:

- **Unigrams and bigrams** of tokens in tweet; we used a Python version of `Tokenizer`, `tokenizer` for Twitter data (Gimpel et al. [2], Myle Ott, 2013 [3]).
- **Tweet metadata**; hashtags, at-mentions and URLs; we parsed the tweet, crawled URLs and used the referred webpage title as supplementary information.

- **Part-of-speech (POS) tags** (bigram and unigram); we used a twitter-specific tagset and tagger presented by Gimpel et al. [2].

Tokenization Splitting a tweet to tokens (separated ordered words) is hard due to irregular punctuation patterns and use of punctuation marks for emoticons. For example, the correct tokenization of 'hello (#hashtag)' is to the four tokens [hello , (, #hashtag ,)], but 'hello (: ' should be split only to the tuple [hello , (:].

Twokenizer is a tokenizer designed for Twitter text in English that addresses these issues. It was originally developed in Python by O'Connor et al., 2010 [1], then improved and ported to Java by Gimpel et al., 2011 [2] and later ported back to Python [3]. We use the last version by Myle Ott.

Metadata Extraction TODO: Gal

Twitter POS Tagging TODO: Gal [2]

3.2 Results

TODO: Daniel?

4 Sentiment Analysis of Tweets

TODO:

5 Named-Entity Recognition in Tweets

TODO:

6 Experimenting with Recent Tweets

TODO: (Gal) Complete this section

keywords Twitter's Search API

7 Conclusions

Future work TODO

References

- [1] BRENDAN O'CONNOR, M. K., AND AHN, D. Tweetmotif: Exploratory search and topic summarization for twitter. In *Proc. ICWSM* (2010).
- [2] KEVIN GIMPEL, NATHAN SCHNEIDER, B. O. D. D. M. J. E. M. H. D. Y. J. F., AND SMITH, N. A. Part-of-speech tagging for twitter: Annotation, features, and experiments. In *Proc. ACL* (2011).
- [3] OTT, M. Python port of Twokenizer. <https://github.com/myleott/ark-twokenize-py>, 2013.

Notes

¹"Disasters on social media" Dataset by CrowdFlower: Contributors looked at over 10,000 tweets culled with a variety of searches like "ablaze", "quarantine", and "pandemonium", then noted whether the tweet referred to a disaster. <https://www.crowdflower.com/wp-content/uploads/2016/03/socialmedia-disaster-tweets-DFE.csv>